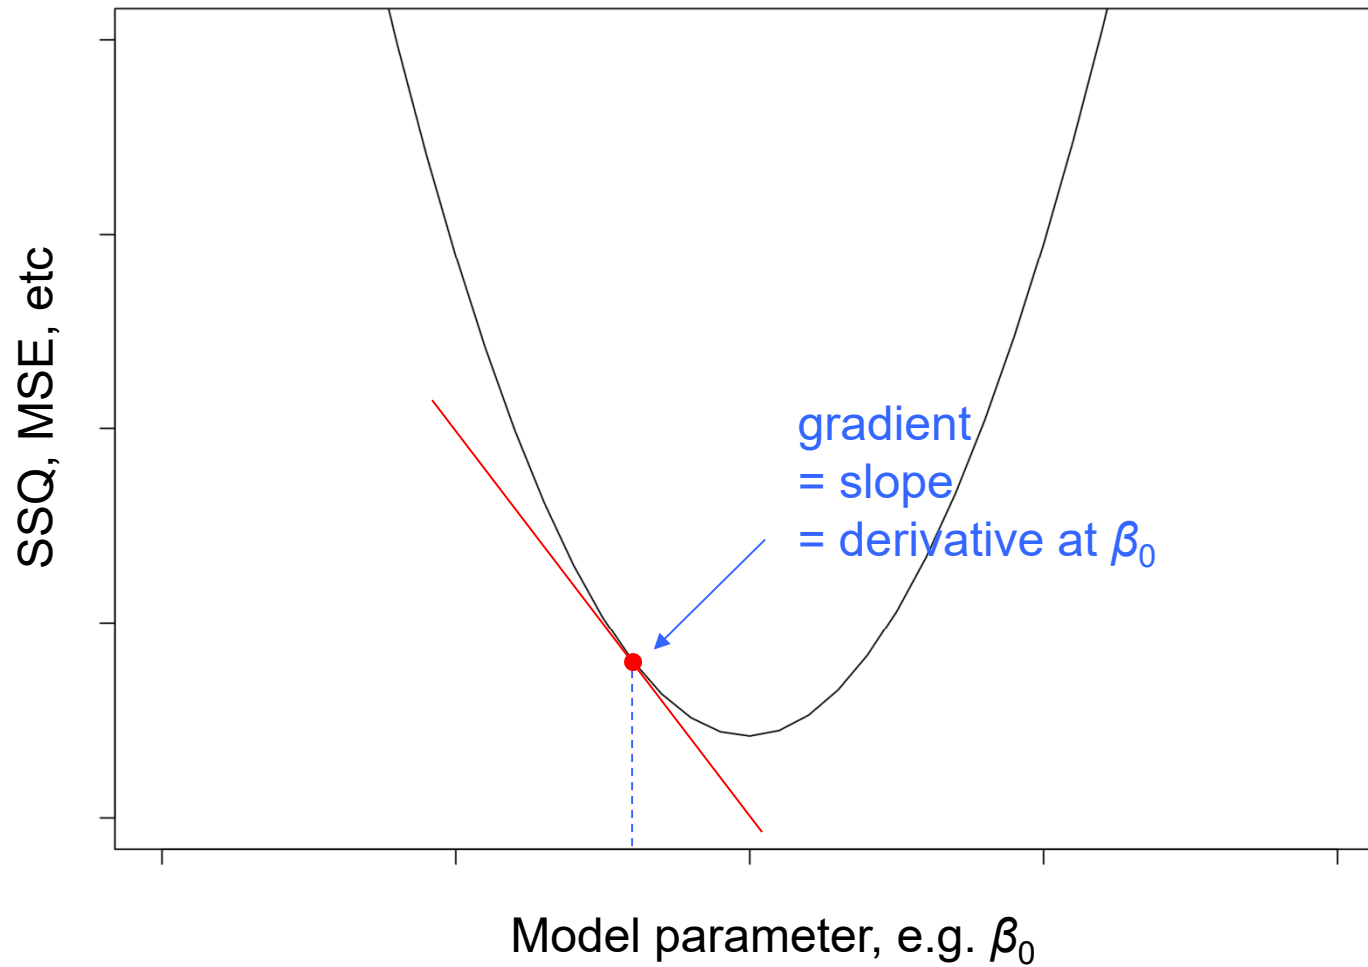


# Today

- Ensemble methods
  - Bagging
  - Random forest
  - Boosting
- But first
  - basic gradient descent
  - boosting is a variant

# Gradient descent



# Finding the gradient

- It turns out that the gradient for a linear model is a function of the residuals
- See math

$$SSQ = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$= \sum_{i=1}^n r_i^2$$

$$= \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

$$= \sum_{i=1}^n y_i^2 - 2y_i \beta_0 - 2y_i \beta_1 x_i + \beta_0^2 + 2\beta_0 \beta_1 x_i + \beta_1^2 x_i^2$$

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$r_i = y_i - \beta_0 - \beta_1 x_i$$

$$\frac{\partial SSQ}{\partial \beta_1} = \sum_{i=1}^n (-2y_i x_i + 2\beta_0 x_i + 2\beta_1 x_i^2)$$

$$= \sum_{i=1}^n -2x_i (y_i - \beta_0 - \beta_1 x_i)$$

$$= -2 \sum_{i=1}^n r_i x_i$$

$$\frac{\partial SSQ}{\partial \beta_0} = \sum_{i=1}^n (-2y_i + 2\beta_0 + 2\beta_1 x_i)$$

$$= -2 \sum_{i=1}^n r_i$$

# Gradient descent

Gradient descent training algorithm for a linear model

```
set lambda (learning rate)
make initial guess for  $\beta_0, \beta_1$ 
for many iterations
    find gradient at  $\beta_0, \beta_1$ 
    step down:  $\beta = \beta - \text{lambda} * \text{gradient}(\beta)$ 
print  $\beta_0, \beta_1$ 
```

# Gradient boosting

## Gradient boosting algorithm (intuitive version)

set lambda (learning rate)

fit a model,  $m$ , to the data

keep a fraction of the model,  $\lambda m$

calculate the left over variation,  $\hat{r}$

repeat until no more systematic variation in  $\hat{r}$

fit a model,  $m$ , to  $\hat{r}$

keep a fraction of the model,  $\lambda m$

add to previous fraction

calculate the left over variation,  $\hat{r}$

# Gradient boosting

## Algorithm

load  $y, x, x_{\text{new}}$

guess parameters:

set  $\hat{f}(x_{\text{new}}) = 0$

set  $r \leftarrow y$  (residuals equal to the data)

for  $m$  in 1 to  $n_{\text{iterations}}$

train model on  $r$  and  $x$

predict residuals,  $\hat{r}_m(x)$ , from trained model

update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$

predict  $y$  increment,  $\hat{f}_m(x_{\text{new}})$ , from trained model

update prediction:  $\hat{f}(x_{\text{new}}) \leftarrow \hat{f}(x_{\text{new}}) + \lambda \hat{f}_m(x_{\text{new}})$

return  $\hat{f}(x_{\text{new}})$

Can be any  
model



Gradient  
descent



# Gradient descent

predict residuals,  $\hat{r}_m(x)$ , from trained model

update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$

**Loss function** is MSE and we are descending its surface


**Gradient** at  $x$  is proportional to the predicted residual  $\hat{r}_m(x)$

$\lambda \hat{r}_m(x)$  is the **increment** taken down the gradient

$r$  gets closer to 0 at each step, so MSE goes down



Boosted linear model (black line) compared to linear regression (red dashed line). They are the same.



# Boosted regression tree

## Algorithm

load  $y, x, x_{\text{new}}$

set parameters:  $\text{tree\_complexity}$ ,  $\text{ntrees}$ ,  $\lambda$

set  $\hat{f}(x_{\text{new}}) = 0$

set  $r \leftarrow y$  (residuals equal to the data)

for  $m$  in 1 to  $\text{ntrees}$

    train  $\text{tree model}$ ,  $m$ , on  $r$  and  $x$


    predict residuals,  $\hat{r}_m(x)$ , from  $\text{trained tree}$

    update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$

    predict  $y$  increment,  $\hat{f}_m(x_{\text{new}})$ , from  $\text{trained tree}$

    update prediction:  $\hat{f}(x_{\text{new}}) \leftarrow \hat{f}(x_{\text{new}}) + \lambda \hat{f}_m(x_{\text{new}})$

return  $\hat{f}(x_{\text{new}})$



# Gradient descent

